Lightweight Object Detection for Driver Monitoring: A Subject-Disjoint Benchmark

Oumar Mamoun Ibrahim
Department of Computer Engineering
University of Sharjah, Sharjah, United Arab Emirates
U22200741@sharjah.ac.ae

Mohamad Khairi bin Ishak
Department of Computer Engineering
University of Sharjah, Sharjah, United Arab Emirates
mishak@sharjah.ac.ae

Khalid Ammar
Department of Electrical and Computer Engineering
Ajman University, Ajman, United Arab Emirates
k.ammar@ajman.ac.ae

Nada Masood Mirza
College of Engineering, Engineering Requirements Unit
United Arab Emirates University, Al Ain, United Arab Emirates
nada.mirza@uaeu.ac.ae

***Abstract*—Reliable driver monitoring requires detectors that generalize to unseen drivers while remaining efficient. This study benchmarks three lightweight object detectors — You Only Look Once (YOLO) 11n, YOLO26n, and nano Fine-grained Distribution Refinement (D-FINE-N) — for visible driver-cue localization under a subject-disjoint protocol. A red–green–blue (RGB) benchmark of 15,723 Driver Monitoring Dataset frames from 14 drivers, including 12,722 negatives, is evaluated across three paired seeds with validation-only operating-point selection and a protected test set. YOLO11n achieves the highest coarse-IoU accuracy and throughput, D-FINE-N the highest calibrated-threshold detection performance, and YOLO26n the strongest strict-IoU localization. Variation across held-out drivers exceeds seed variation, underscoring the importance of subject-level reporting. A complementary eight-system near-infrared study confirms that greater negative training exposure does not consistently improve performance. Detector selection should therefore reflect deployment-specific trade-offs.**

***Index Terms*—driver monitoring, object detection, subject-disjoint evaluation, visible driver cues, near-infrared imaging, real-time inference**

I. Introduction

Road traffic injuries claim approximately 1.19 million lives annually and impose a substantial economic burden worldwide [1]. In response, Euro NCAP now awards safety-assist credit for camera-based driver monitoring [2], accelerating the deployment of in-cabin driver monitoring systems (DMSs). These systems must operate robustly under changes in driver identity, illumination, viewpoint, and occlusion. Whole-frame recognition can describe an activity, but it may also exploit identity or cabin context without showing the visual evidence used for a decision. Object detection provides a complementary formulation: it localizes cue-specific evidence, makes false positives and misses spatially interpretable, and measures detections on frames containing no target cue. A single-frame detector also avoids the latency and architectural complexity of a temporal model when real-time throughput is the experimental focus. This formulation does *not* infer fatigue, intent, duration, or another temporal driver state.

DMD provides multimodal attention and alertness recordings [3], while Drive&Act contains synchronized color, NIR, depth, and pose streams for fine-grained activities [4]. Both contain temporally correlated video and repeated observations of each driver.

Evaluation can therefore be optimistic if neighboring frames or identities cross partitions, natural negative frames are removed, the test set determines the confidence threshold, or latency boundaries differ. Existing driver-detection studies on DMD [5] and posture-classification benchmarks such as the State Farm challenge [6] further illustrate the diversity of evaluation choices and the difficulty of cross-paper comparison.

Recent work proposes specialized lightweight networks for in-cabin recognition and detection [7]–[10]. The objective of this study is to investigate how lightweight object detectors generalize to unseen drivers and how their localization performance, detection reliability, false-positive behavior, and computational requirements differ under a controlled subject-disjoint evaluation. The study considers complete, publicly released detector systems without task-specific architectural modification. A separate exploratory NIR experiment examines whether increased exposure to negative training examples produces consistent changes in detector performance across different model families. The contributions are:

- a primary subject-disjoint RGB benchmark for four explicitly defined visible cues, with all sampled frames preserved and annotation structure automatically audited;
- validation-only operating-point selection followed by protected testing, three paired RGB seeds, class- and subject-level analysis, false detections on negative frames, and synchronized latency; and
- a checksum-backed reproduction workflow plus a clearly separated, one-seed, eight-system NIR training-negative exposure study with deterministic 1-FPS midpoint sampling.

II. Related Work

A. Driver Recognition and Cue Localization

MobileVGG and convolution–transformer hybrids demonstrate efficient distracted-driver classification [7], [11]. POGUISE+ uses pose and object guidance for efficient driver-action recognition in Drive&Act NIR imagery [12]. Such whole-frame or temporal recognition tasks are not interchangeable with the present bounding-box task. Here, a correct prediction must identify both a cue class and its spatial evidence in one frame; no temporal state is estimated.

Detection-specific work modifies YOLO backbones, attention, or multiscale aggregation for dangerous-driving or fatigue cues [8]–[10]. Comparative studies assess convolutional neural networks against transformers [13] or several YOLO generations [14]. The YOLO family originates with single-shot grid regression [15] and has evolved through multiple architectural revisions [16]–[18]. However, differences in driver separation, retained negatives, operating-point calibration, and runtime protocol limit direct cross-paper ranking.

B. Released Detector Systems

The Detection Transformer (DETR) casts object detection as bipartite set prediction [19]; the Real-Time Detection Transformer (RT-DETR) makes this formulation practical for real-time use [20]; RT-DETRv2 extends it with a bag-of-freebies baseline [21]; and D-FINE refines box-edge distributions for improved localization [22]. The RGB and NIR comparisons use D-FINE-N, YOLO11n [17], and YOLO26n [18] as shipped systems. The NIR study additionally evaluates YOLOv8n [23], YOLOX-Nano [24], YOLOv10n [25], SSDLite-MobileNetV3-Large [26], and RT-DETRv2-S [21]. Native training and inference recipes remain part of each system: the pinned YOLO26 and D-FINE paths are end-to-end, whereas YOLO11n retains non-maximum suppression (NMS) [27]. Recent edge-deployment work demonstrates that such lightweight systems can approach real-time performance on low-cost hardware [28]. This is therefore a system-level benchmark, not an isolation of architecture alone.

Across these lines of work, no prior study combines subject-disjoint splitting, validation-only threshold selection with a protected test set, retention of all negative frames, and synchronized latency profiling in a single controlled comparison of shipped lightweight detectors. The present benchmark addresses this gap.

III. Methodology

A. Task Definition and Workflow

The task is bounding-box localization of four visible RGB cues: yawning, hand over mouth, drinking, and phone use. Detection is used because it spatially localizes the evidence, discourages whole-frame identity and cabin shortcuts, and exposes errors on negative frames. Figure 1 summarizes the shared protocol. Dataset construction, splitting, training, and validation precede a checksum and confirmation gate; no test-derived threshold or model change is permitted. The RGB benchmark is primary. NIR is a distinct secondary experiment rather than a paired modality comparison.

TABLE I
Frozen RGB Box Semantics and Test Support

Cue	Bounding-box target	Test n	Median area
Yawning	Mouth	33	1.38%
Hand over mouth	Full face and covering hand	28	24.67%
Drinking	Visible hand/container–mouth interaction	56	16.09%
Phone use	Visible hand–phone interaction	497	8.37%

B. RGB Dataset and Annotation Protocol

The RGB set comprises 15,723 face crops sampled at 1 FPS from 81 DMD recordings of 14 drivers. A single annotator completed the RGB annotations in one annotation pass. No independent second annotator was available for semantic agreement assessment. All frames were retained without post-sampling curation: 3,001 positive and 12,722 negative. Table I fixes the spatial ontology and cue-specific box semantics.

Each frame contains zero or one box. Yawning and hand-over-mouth are mutually exclusive; when a hand covers a yawn, hand-over-mouth takes precedence. Boxes enclose only the visible target. Occluded, truncated, or out-of-frame anatomy is never estimated, although a partially occluded cue remains valid when visually identifiable. No arbitrary minimum-pixel cutoff is applied at 640×640.

An automated COCO audit found 15,723 unique images and 3,001 unique annotations, with zero duplicate identifiers or file names, orphan annotations, invalid categories, nonfinite or nonpositive boxes, out-of-bounds boxes, or multi-annotation frames. The smallest annotated area was 0.81% of the frame; 139 of 159 yawning boxes occupied below 2%. These checks verify structural consistency of the annotation files but do not assess semantic agreement or labeling accuracy.

C. Split, Models, and Training

An exhaustive 8:3:3 subject assignment minimizes the worst relative class-distribution deviation, with root-mean-square deviation as a tie-breaker. The frozen training, validation, and test partitions contain 9,087, 3,423, and 3,213 frames. The three test subjects contribute 1,080, 920, and 1,213 frames and 178, 183, and 253 positives, respectively.

COCO-pretrained YOLO11n [17], YOLO26n [18], and D-FINE-N [22] use 640×640 input, physical batch size 8, four-step gradient accumulation for an effective batch of 32, 220 epochs, automatic mixed precision (AMP), and no early stopping. A fixed patch normalizes an incomplete final accumulation window by its actual sample count. Native recipes are retained: Ultralytics `optimizer=auto` with a linear schedule, and the official D-FINE AdamW/MultiStepLR recipe. Seeds $\{13, 37, 73\}$ were declared in advance, and no run was selected for presentation. Table II consolidates the shared settings.

D. Validation, Test Metrics, and Profiling

The best model-native validation checkpoint is retained. Its confidence threshold τ^* maximizes validation micro-F_1 over

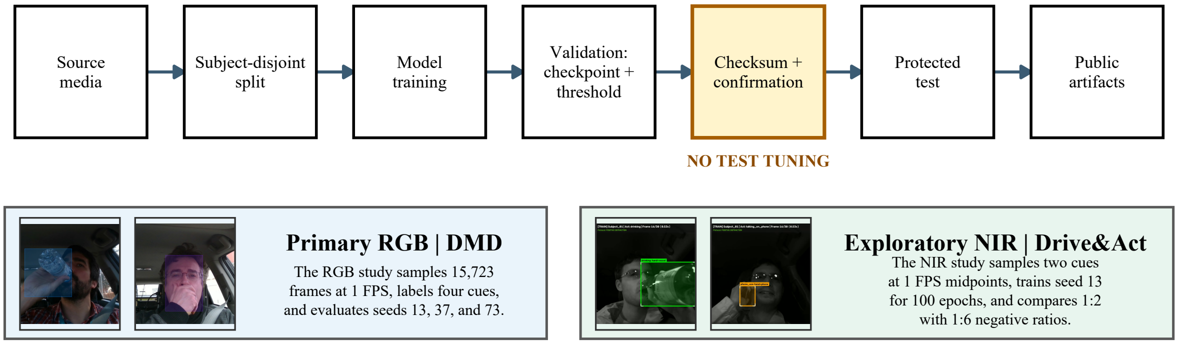

Fig. 1. End-to-end benchmark workflow with representative annotated DMD and Drive&Act frames. RGB and NIR share the validation-before-test gate but retain dataset-specific sampling, ontology, seeds, and training conditions. FPS denotes frames per second.

TABLE II
SHARED RGB TRAINING CONFIGURATION

Setting	Value
Input resolution	640 × 640
Physical / effective batch	8 / 32 (4-step accumulation)
Epochs	220 (no early stopping)
Precision	AMP (FP16/FP32)
Pretrained backbone	COCO
Seeds	$\{13, 37, 73\}$ (predeclared)
YOLO optimizer	`auto` + linear schedule
D-FINE optimizer	AdamW + MultiStepLR

$[0.01, 0.99]$, with higher precision and then higher τ as tie-breakers. At an intersection over union (IoU) of 0.5, predictions are greedily matched one-to-one within class. Protected testing reports COCO mAP_{50} and $\mathrm{mAP}_{50:95}$ [29], precision, recall, micro- and macro-F_1, and

$$\mathrm{FD}_{100} = 100\frac{\mathrm{FP}_{\mathrm{neg}}}{N_{\mathrm{neg}}}, \tag{1}$$

where $\mathrm{FP}_{\mathrm{neg}}$ is the number of above-threshold false-positive detections on negative frames. We define FD_{100} (false-detection rate per 100 negative frames) as a simple diagnostic ratio — distinct from the false discovery rate $\mathrm{FP}/(\mathrm{FP}+\mathrm{TP})$ used in hypothesis testing — that measures how often the detector fires on frames where no target cue is present. It is a static per-frame count, not a modeled nuisance-alert rate; no temporal confirmation or debouncing is applied.

RGB results are the sample mean and sample standard deviation (SD) over three seeds. Because the subject split is fixed, SD measures optimization variability, not generalization uncertainty across held-out drivers. Each test subject is therefore reported separately, and paired per-seed differences are checked for sign consistency. No significance test is claimed for $n = 3$.

For all evaluated systems, model selection and operating-point determination were performed exclusively on the validation set, while the held-out test partition remained fixed and was used only for final evaluation.

Batch-1 16-bit floating-point (FP16) profiling uses an NVIDIA GeForce RTX 4060 8-GB GPU. CUDA synchronization brackets model-forward and tensor-to-final-detection intervals. Disk input/output, decoding, preprocessing, and metric computation are excluded. We report p50 tensor-to-final latency and sustained FPS; these measurements do not represent CPU or embedded-device deployment.

E. Exploratory NIR Protocol

A practical question for DMS deployment is whether exposing a detector to more negative examples during training consistently reduces false detections. Because this question cannot be answered from the single-condition RGB benchmark above, a separate NIR experiment is introduced using a different dataset and a broader set of detector families. The results are not comparable to the RGB track; they address a complementary training-design question.

The secondary NIR source pool contains 3,786 one-second front-top active-NIR (850 nm) Drive&Act snippets from 15 drivers. Twenty-three unilluminated task snippets are excluded before splitting, leaving 3,763. Drinking and phone-use tracklets provide two positive classes. Each snippet contributes one deterministic frame at 1 FPS: the temporal midpoint at $t = 0.5$ s, or zero-based source-frame offset 14. At that instant, the normalized box coordinates (x, y, width, and height) are linearly interpolated between the nearest human-labeled tracklet keyframes that bracket the midpoint; when no bracketing pair exists, the nearest endpoint box is used. The driver crop is then resized to 640 × 640. Unlike RGB, NIR therefore applies dataset-specific frame exclusion and derives midpoint annotations from human tracklets.

A fixed 9:3:3 subject partition is shared across all sixteen configured model–condition runs. The eight systems span one-stage convolutional, SSD, and transformer-based detector families. Validation and test contain identical ordered examples. Training always contains 270 positive snippets and either 540 or 1,620 negative snippets, corresponding to nominal positive:negative ratios of 1:2 and 1:6. The 1:2 negatives are a subject-stratified seed-13 subset of the 1:6 pool.

All runs use physical batch 8, four-step accumulation, a fixed 100-epoch maximum without early stopping, and the best model-native validation checkpoint. Native optimizers and schedules are retained. The D-FINE augmentation transition is placed at epoch 67. The 1:2 and 1:6 conditions produce 2,600 and 6,000 optimizer updates, respectively. Thus, 1:6 has three times the unique negatives and approximately 2.31 times the optimization updates. It is consequently a *training-negative exposure study*, not a causal class-ratio ablation under matched training signal.

All NIR jobs use seed 13, and operating metrics follow the RGB protocol. Since one frame represents each snippet, no temporal aggregation is performed. Micro-F_1 and FD_{100} retain the detection-level definitions above; macro-F_1 uses the stored class-presence localization metric, in which repeated same-class detections on a frame count once.

IV. Results and Analysis

A. Aggregate RGB Results

Table III summarizes all nine RGB runs. YOLO11n leads mAP_{50} and throughput, YOLO26n leads $\mathrm{mAP}_{50:95}$, and D-FINE-N leads micro- and macro-F_1. The D-FINE advantage is concentrated at the calibrated operating point rather than strict-IoU AP and costs 8.0–26.4 FPS relative to the YOLO systems. The evidence supports an accuracy–efficiency trade-off, not universal superiority.

B. Class Imbalance and Paired-Seed Evidence

Phone use accounts for 497 of 614 positive test instances, whereas hand over mouth, yawning, and drinking have only 28, 33, and 56 instances. This imbalance can influence learned representations, validation-selected thresholds, and the pooled micro-F_1. It does not directly dominate COCO mAP, because AP is first computed by category and then averaged equally across the four categories.

The class-wise mean differences, YOLO26n minus YOLO11n, are $+6.21$ points for hand over mouth, -3.85 for yawning, $+1.02$ for drinking, and $+3.02$ for phone use. D-FINE-N obtains the highest mean AP for yawning, drinking, and phone use, but the lowest for hand over mouth. Reporting macro-F_1 and per-class AP beside micro-F_1 prevents the majority phone-use class from silently defining the conclusion, although the wide SD for low-support cues still warrants caution. All three seeds agree on the direction of the YOLO26n $\mathrm{mAP}_{50:95}$ advantage over YOLO11n, providing consistent paired evidence despite the small seed count.

C. Held-Out Subject Sensitivity

Averaged across seeds, YOLO11n micro-F_1 is 94.94%, 62.41%, and 94.72% for Subjects 05, 10, and 12; YOLO26n obtains 96.11%, 54.11%, and 92.29%; and D-FINE-N obtains 98.20%, 78.79%, and 94.30%. The between-subject change is far larger than the pooled seed SD. Reporting only the pooled test set would therefore hide the difficult driver.

D. Qualitative Error Analysis

Figure 3 shows predeclared seed-13 examples from frozen predictions for all three RGB detectors. Correct detections localize the intended interaction, while the false-negative examples expose glare and small drinking targets. More importantly, nominal negatives can appear visually cue-like. They may reflect difficult ontology boundaries or missed labels, so the structural annotation audit alone cannot establish semantic validity. The examples are retained because they expose the benchmark's principal annotation risk.

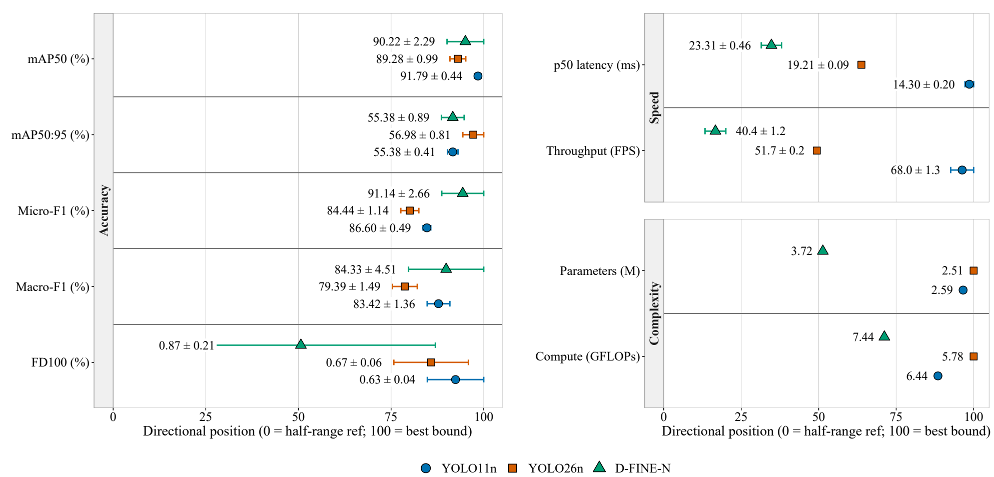

Fig. 2. Directional dot-and-whisker plot for YOLO11n, YOLO26n, and D-FINE-N. Rows are grouped into accuracy (left), speed, and complexity (right); right indicates better performance, and whiskers show sample SD across seeds.

TABLE III
Mean and sample standard deviation across three seeds for the RGB track on the protected test split. Bold underlined values indicate the best result per column.

Model	mAP_{50}	$\mathrm{mAP}_{50:95}$	Micro-F_1	Macro-F_1	FD_{100}	p50 (ms)	FPS	Params (M)	GFLOPs
YOLO11n	**91.79±0.44**	55.38±0.41	86.60±0.49	83.42±1.36	**0.63±0.04**	**14.30±0.20**	**68.0±1.3**	2.59	6.44
YOLO26n	89.28±0.99	**56.98±0.81**	84.44±1.14	79.39±1.49	0.67±0.06	19.21±0.09	51.7±0.2	**2.51**	**5.78**
D-FINE-N	90.22±2.29	55.38±0.89	**91.14±2.66**	**84.33±4.51**	0.87±0.21	23.31±0.46	40.4±1.2	3.72	7.44

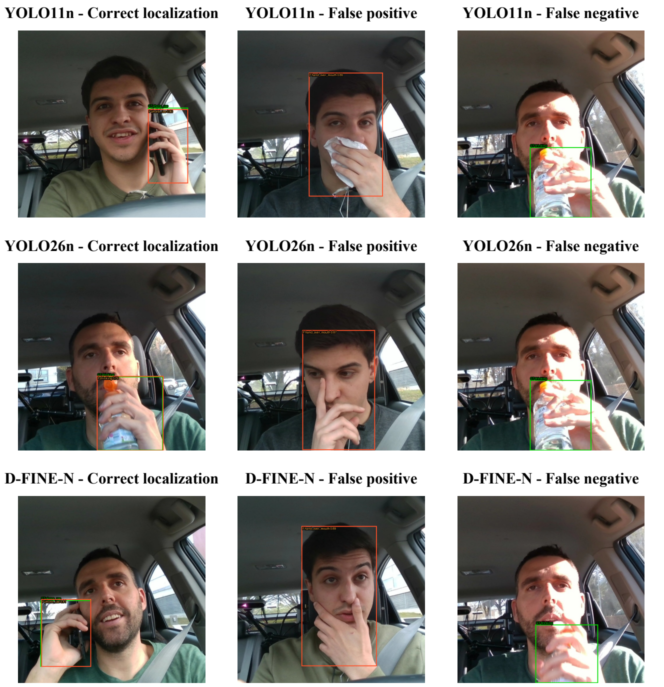

Fig. 3. Representative seed-13 outcomes for YOLO11n, YOLO26n, and D-FINE-N. Green boxes denote ground truth and orange boxes denote predictions at the validation-selected threshold; columns show a correct localization, a false detection on a nominal negative, and a false negative.

E. Exploratory NIR Results

All 16 protected NIR passes are complete. Table IV reports both exposures for YOLO11n, YOLO26n, D-FINE-N, SSDLite-MobileNetV3-Large, RT-DETRv2-S, YOLOX-Nano, YOLOv10n, and YOLOv8n, while Fig. 5 plots the same checksum-backed results. Figure 4 shows agreement-selected qualitative outcomes across the 16 runs.

Increasing negative exposure from 1:2 to 1:6 improves $\mathrm{mAP}_{50:95}$ for only YOLO26n, SSDLite-MobileNetV3-Large, and RT-DETRv2-S; it lowers AP for the other five systems. The changes range from $+5.52$ percentage points for YOLO26n to -7.71 for YOLOX-Nano. Micro-F_1 improves only for SSDLite-MobileNetV3-Large and decreases for the other seven systems. The direction is therefore model-dependent rather than a uniform benefit of 1:6 exposure. RT-DETRv2-S has the highest strict-IoU AP at both ratios, while YOLOv8n has the highest micro-F_1 and lowest tensor-to-final p50 latency at both ratios.

V. Discussion

A. Performance and Deployment Trade-offs

The results reveal clear trade-offs among localization performance, detection reliability, and computational efficiency. In RGB, YOLO11n provides the strongest balance of coarse-IoU accuracy and throughput, YOLO26n achieves the best strict-IoU localization, and D-FINE-N attains the highest F_1 performance at greater inference latency. The NIR results similarly show that performance varies across models and evaluation criteria.

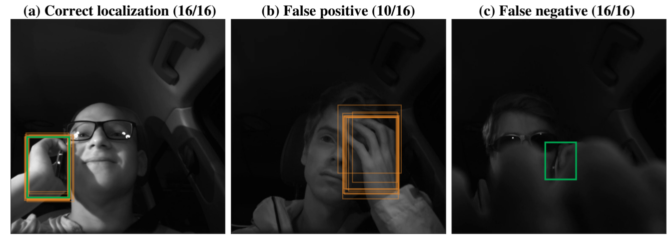

Fig. 4. Agreement-selected NIR outcomes at frozen thresholds. Green denotes ground truth and orange denotes predictions. The correct localization and false negative are unanimous; the most recurrent negative-frame false positive occurs in 10/16 runs.

TABLE IV
Exploratory NIR Test Results at Seed 13. Per-model ratio winners are bold underlined; displayed ties are unmarked.

Model	Neg.	$\mathrm{mAP}_{50:95}$	Micro-F_1	Macro-F_1	FD_{100}	p50
YOLO11n	1:2	**42.62**	**71.91**	61.94	**0.14**	15.22
YOLO11n	1:6	37.07	69.15	**70.32**	0.41	**13.74**
YOLO26n	1:2	35.75	**76.54**	**83.62**	2.17	17.11
YOLO26n	1:6	**41.28**	65.96	65.57	**0.68**	**16.84**
D-FINE-N	1:2	**46.38**	**83.58**	**80.25**	**0.54**	24.58
D-FINE-N	1:6	46.30	79.23	76.45	0.81	**24.10**
SSDLite-MNV3-L	1:2	28.43	73.39	69.97	3.11	**14.76**
SSDLite-MNV3-L	1:6	**30.76**	**78.73**	**73.38**	**1.35**	15.93
RT-DETRv2-S	1:2	48.14	**84.42**	**78.14**	0.27	25.61
RT-DETRv2-S	1:6	**48.51**	79.79	76.80	0.27	**24.66**
YOLOX-Nano	1:2	**33.95**	**79.23**	**76.36**	0.95	**15.46**
YOLOX-Nano	1:6	26.24	67.37	61.31	**0.41**	16.24
YOLOv10n	1:2	**36.67**	**69.23**	**65.40**	2.17	15.41
YOLOv10n	1:6	32.58	66.67	64.32	**2.03**	**13.45**
YOLOv8n	1:2	**39.21**	**84.51**	**83.71**	0.27	11.74
YOLOv8n	1:6	37.77	80.81	76.75	0.27	**10.24**

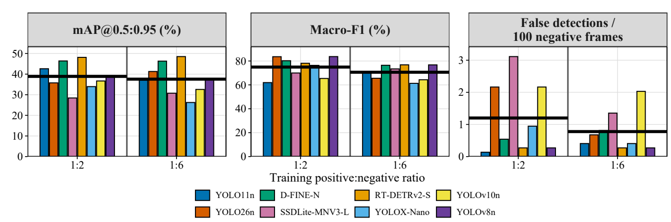

Fig. 5. Sixteen NIR protected-test results at seed 13 and deterministic 1-FPS midpoint sampling. Distinct colors identify the eight models at 1:2 and 1:6 training exposures; the dark-grey rule separates the ratio sections, and each thick black segment marks the eight-model mean for its ratio.

These findings indicate that no single detector is optimal across all requirements, and model selection should therefore reflect the specific accuracy, reliability, and efficiency demands of the intended deployment.

B. Class and Subject Effects

The phone-use majority is a real benchmark characteristic and a limitation for pooled operating metrics, but it does not invalidate YOLO26n's strict-IoU advantage. Equal-weight class AP, macro-F_1, per-class supports, and paired-seed directions separate that result from majority-class dominance. The narrow aggregate margin, low support for three categories, and drinking sign reversal mean the accurate conclusion is a trade-off rather than a definitive rank. The difficult Subject 10 further shows that seed variation on a single split understates deployment variation across people.

C. Interpretation Scope

Subject-disjoint evaluation reduces identity leakage but does not establish population-level generalization from only 14 drivers and three held-out test identities. A single annotator and one pass provide internally consistent data but no independent semantic agreement estimate; the qualitative nominal-negative case makes this concrete. Future annotation rounds should compute inter-annotator agreement (e.g., Cohen's κ on a stratified subset of each cue class) to quantify this risk. Correlated frames remain within RGB recordings, and single-frame detection neither estimates temporal fatigue nor models alert debouncing.

NIR avoids within-snippet duplication through one midpoint frame, but 1-FPS sampling does not characterize motion within the second. NIR additionally differs in dataset, ontology, annotation source, and preprocessing, uses one seed, and changes negative diversity together with training steps. It cannot be interpreted as a controlled RGB-versus-NIR modality effect.

Finally, RTX 4060 timing excludes preprocessing and is not an embedded-hardware measurement. These boundaries define what the benchmark supports without removing its value as a transparent, reproducible comparison.

VI. Conclusion and Future Work

This study presented a subject-disjoint RGB benchmark and a complementary NIR negative-exposure experiment for lightweight driver-cue detection. The central finding is that no single detector dominates all evaluation axes: coarse-IoU accuracy and throughput, strict-IoU localization, and calibrated-threshold detection reliability each favor a different system.

For practitioners, the results suggest concrete selection criteria. When false alarms are costlier than missed cues — for example, nuisance alerts that erode driver trust — YOLO11n's lower FD_{100} and highest throughput make it a natural starting point. When precise spatial localization matters, such as feeding downstream gaze or posture models, YOLO26n's strict-IoU advantage is more relevant. When the deployment can absorb additional latency in exchange for stronger F_1 at the calibrated threshold, D-FINE-N offers the highest detection reliability. The large performance gap on the difficult held-out Subject 10 further underscores that aggregate metrics can mask deployment-critical failure modes; subject-level reporting should be standard practice.

Future work should expand the number and demographic diversity of held-out drivers, incorporate independent annotation review to quantify semantic agreement, repeat the analysis across multiple subject partitions to separate split-specific effects, and evaluate complete inference pipelines on embedded hardware.

Data and code: Authored annotations, frozen partitions and checksums, sanitized predictions, evaluation tools, and publication artifacts are available at https://github.com/IamOumarIbrahim/lightweight-driver-behavior-detection; licensed DMD and Drive&Act media and trained checkpoints remain local under their original access terms.

References

[1] World Health Organization, "Global status report on road safety 2023," World Health Organization, Geneva, Tech. Rep., 2023, licence: CC BY-NC-SA 3.0 IGO. [Online]. Available: https://www.who.int/publications/i/item/9789240086517

[2] Euro NCAP, "Assessment protocol – safety assist: Driver monitoring," European New Car Assessment Programme, Tech. Rep., 2025, v10.0.3, Implementation 2025. [Online]. Available: https://www.euroncap.com/en/for-engineers/protocols/safety-assist/

[3] J. D. Ortega *et al.*, "DMD: A large-scale multi-modal driver monitoring dataset for attention and alertness analysis," in *Computer Vision – ECCV 2020 Workshops*. Springer, 2020, pp. 387–405.

[4] M. Martin *et al.*, "Drive&Act: A multi-modal dataset for fine-grained driver behavior recognition in autonomous vehicles," in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019, pp. 2801–2810.

[5] R. Jain, A. Ballav, S. Haresh, G. Pooja, and H. Gupta, "Driver distraction detection using YOLO variants on the DMD dataset," in *2024 International Conference on Advances in Computing, Communication and Applied Informatics*, 2024, pp. 1–6.

[6] Y. Abouelnaga, H. M. Eraqi, and M. N. Moustafa, "Real-time distracted driver posture classification," in *Machine Learning for Intelligent Transportation Systems Workshop, NeurIPS 2018*, 2018.

[7] B. Baheti, S. Talbar, and S. Gajre, "Towards computationally efficient and realtime distracted driver detection with MobileVGG network," *IEEE Transactions on Intelligent Vehicles*, vol. 5, no. 4, pp. 565–574, 2020.

[8] C. Wang, M. Lin, L. Shao, and J. Xiang, "MF-YOLO: A lightweight method for real-time dangerous driving behavior detection," *IEEE Transactions on Instrumentation and Measurement*, vol. 73, pp. 1–13, 2024.

[9] R. Li *et al.*, "YOLO-SGC: A dangerous driving behavior detection method with multiscale spatial-channel feature aggregation," *IEEE Sensors Journal*, vol. 24, no. 21, pp. 36 044–36 056, 2024.

[10] Z. Jin and L. Dong, "YOLO11-CR: A lightweight convolution-and-attention framework for accurate fatigue driving detection," 2025, arXiv:2508.13205.

[11] Z. Li, X. Zhao, F. Wu, D. Chen, and C. Wang, "A lightweight and efficient distracted driver detection model fusing convolutional neural network and vision transformer," *IEEE Transactions on Intelligent Transportation Systems*, vol. 25, no. 12, pp. 19 962–19 978, 2024.

[12] R. Pizarro, R. Valle, R. Barea, J. M. Buenaposada, L. Baumela, and L. M. Bergasa, "PO-GUISE+: Pose and object guided transformer token selection for efficient driver action recognition," *IEEE Transactions on Intelligent Transportation Systems*, vol. 27, no. 7, pp. 8424–8436, 2026.

[13] H. V. Koay, J. H. Chuah, and C.-O. Chow, "Convolutional neural network or vision transformer? benchmarking various machine learning models for distracted driver detection," in *2021 IEEE Region 10 Conference (TENCON)*, 2021, pp. 417–421.

[14] D. Herath, C. Abeyrathne, and P. Jayaweera, "Vision-based driver drowsiness monitoring: Comparative analysis of YOLOv5–v11 models," 2025, arXiv:2509.17498.

[15] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You only look once: Unified, real-time object detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 779–788.

[16] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, "YOLOv4: Optimal speed and accuracy of object detection," 2020, arXiv:2004.10934.

[17] G. Jocher and J. Qiu, "Ultralytics YOLO11," https://docs.ultralytics.com/models/yolo11/, 2024, version 11.0. Accessed: Aug. 25, 2026.

[18] Ultralytics, "YOLO26: End-to-end NMS-free detection," https://docs.ultralytics.com/models/yolo26/, 2026, accessed: Aug. 25, 2026.

[19] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko, "End-to-end object detection with transformers," in *Computer Vision – ECCV 2020*. Springer, 2020, pp. 213–229.

[20] Y. Zhao *et al.*, "DETRs beat YOLOs on real-time object detection," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 16 965–16 974.

[21] W. Lv, Y. Zhao, Q. Chang, K. Huang, G. Wang, and Y. Liu, "RT-DETRv2: Improved baseline with bag-of-freebies for real-time detection transformer," 2024, arXiv:2407.17140. [Online]. Available: https://arxiv.org/abs/2407.17140

[22] Y. Peng, H. Li, P. Wu, Y. Zhang, X. Sun, and F. Wu, "D-FINE: Redefine regression task of DETRs as fine-grained distribution refinement," in *The Thirteenth International Conference on Learning Representations*, 2025. [Online]. Available: https://openreview.net/forum?id=q4GJBvGsMm

[23] G. Jocher, A. Chaurasia, and J. Qiu, "Ultralytics YOLOv8," https://github.com/ultralytics/ultralytics, 2023, version 8.0. Accessed: Aug. 25, 2026.

[24] Z. Ge, S. Liu, F. Wang, Z. Li, and J. Sun, "YOLOX: Exceeding YOLO series in 2021," 2021, arXiv:2107.08430.

[25] A. Wang *et al.*, "YOLOv10: Real-time end-to-end object detection," in *Advances in Neural Information Processing Systems*, 2024.

[26] A. Howard *et al.*, "Searching for MobileNetV3," in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019, pp. 1314–1324.

[27] N. Bodla, B. Singh, R. Chellappa, and L. S. Davis, "Soft-NMS – improving object detection with one line of code," in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 5561–5569.

[28] V. Ahsani, B. H. Khalaj, and H. Shah-Mansouri, "Real-time in-cabin driver behavior recognition on low-cost edge hardware," 2025, arXiv:2512.22298.

[29] T.-Y. Lin *et al.*, "Microsoft COCO: Common objects in context," in *Computer Vision – ECCV 2014*. Springer, 2014, pp. 740–755.